

MICRO-DEGREE

Künstliche Intelligenz und Gesellschaft

***Artificial Intelligence
and Society***



Micro-Degree Artificial Intelligence and Society

– Technical Aspects of AI –

Lecture 4.3: Generative Language Models

Jana Lasser

Research group

Complex Social & Computational Systems (CS)²

IDEa_Lab


Das Interdisziplinäre Digitale
Labor der Universität Graz



Generative Language Models

- Generative language models – sometimes also "Large Language Models" or "LLMs" – can produce coherent sounding text on demand.
- They have become extremely widespread since the release of the chat bot ChatGPT in 2022.
- To train generative language models, all types of machine learning are needed.

Step 1: Pre-training with Self-Supervised Learning



**giant
text corpus**


GPT-3's training corpus includes:

- **CommonCrawl** (400 billion tokens, 46% English) – low quality.
- Reddit (19 billion tokens), Wikipedia (3 billion tokens) – high quality.
- Books (67 billion tokens) – high quality.

See also Brown et al., [Language Models are Few-Shot Learners](#), *arXiv* (2020).

**legal aspects
of AI**

Step 1: Pre-training with Self-Supervised Learning



GPT-3's training corpus includes:

- **CommonCrawl** (400 billion tokens, 46% English) – low quality.
- Reddit (19 billion tokens), Wikipedia (3 billion tokens) – high quality.
- Books (67 billion tokens) – high quality.

See also Brown et al., [Language Models are Few-Shot Learners](#), *arXiv* (2020).

Cost of training GPT-4:
\$100 mio


**economic
aspects of AI**

Pre-trained Language Model

- A language model pre-trained in this way is already very good at generating **coherent sequences of words**.
- It "speaks" the languages sufficiently present in the training dataset and has learned syntax and grammar structures.
- Pre-trained language models are **very large**: they are neural networks with up to 671 billion parameters* (weights). This corresponds to about 1.5 TB of hard drive space just for storing the model weights.
- However, the model cannot yet **follow instructions**. Produced text follows the structure of books or articles rather than conversations.

*Size of the most recent [DeepSeek model V3](#).


Step 2: Supervised Learning for Instruction Following




Instruction-Tuned Language Model

- The "fine-tuning" of the model with a small dataset of instruction-response pairs changes the model's weights only superficially – the basic structure of the language remains intact.
- An instruction-tuned language model can respond meaningfully to instructions as an input sequence.
- However, it can still produce answers that may be **unfriendly, dangerous, or misleading**.

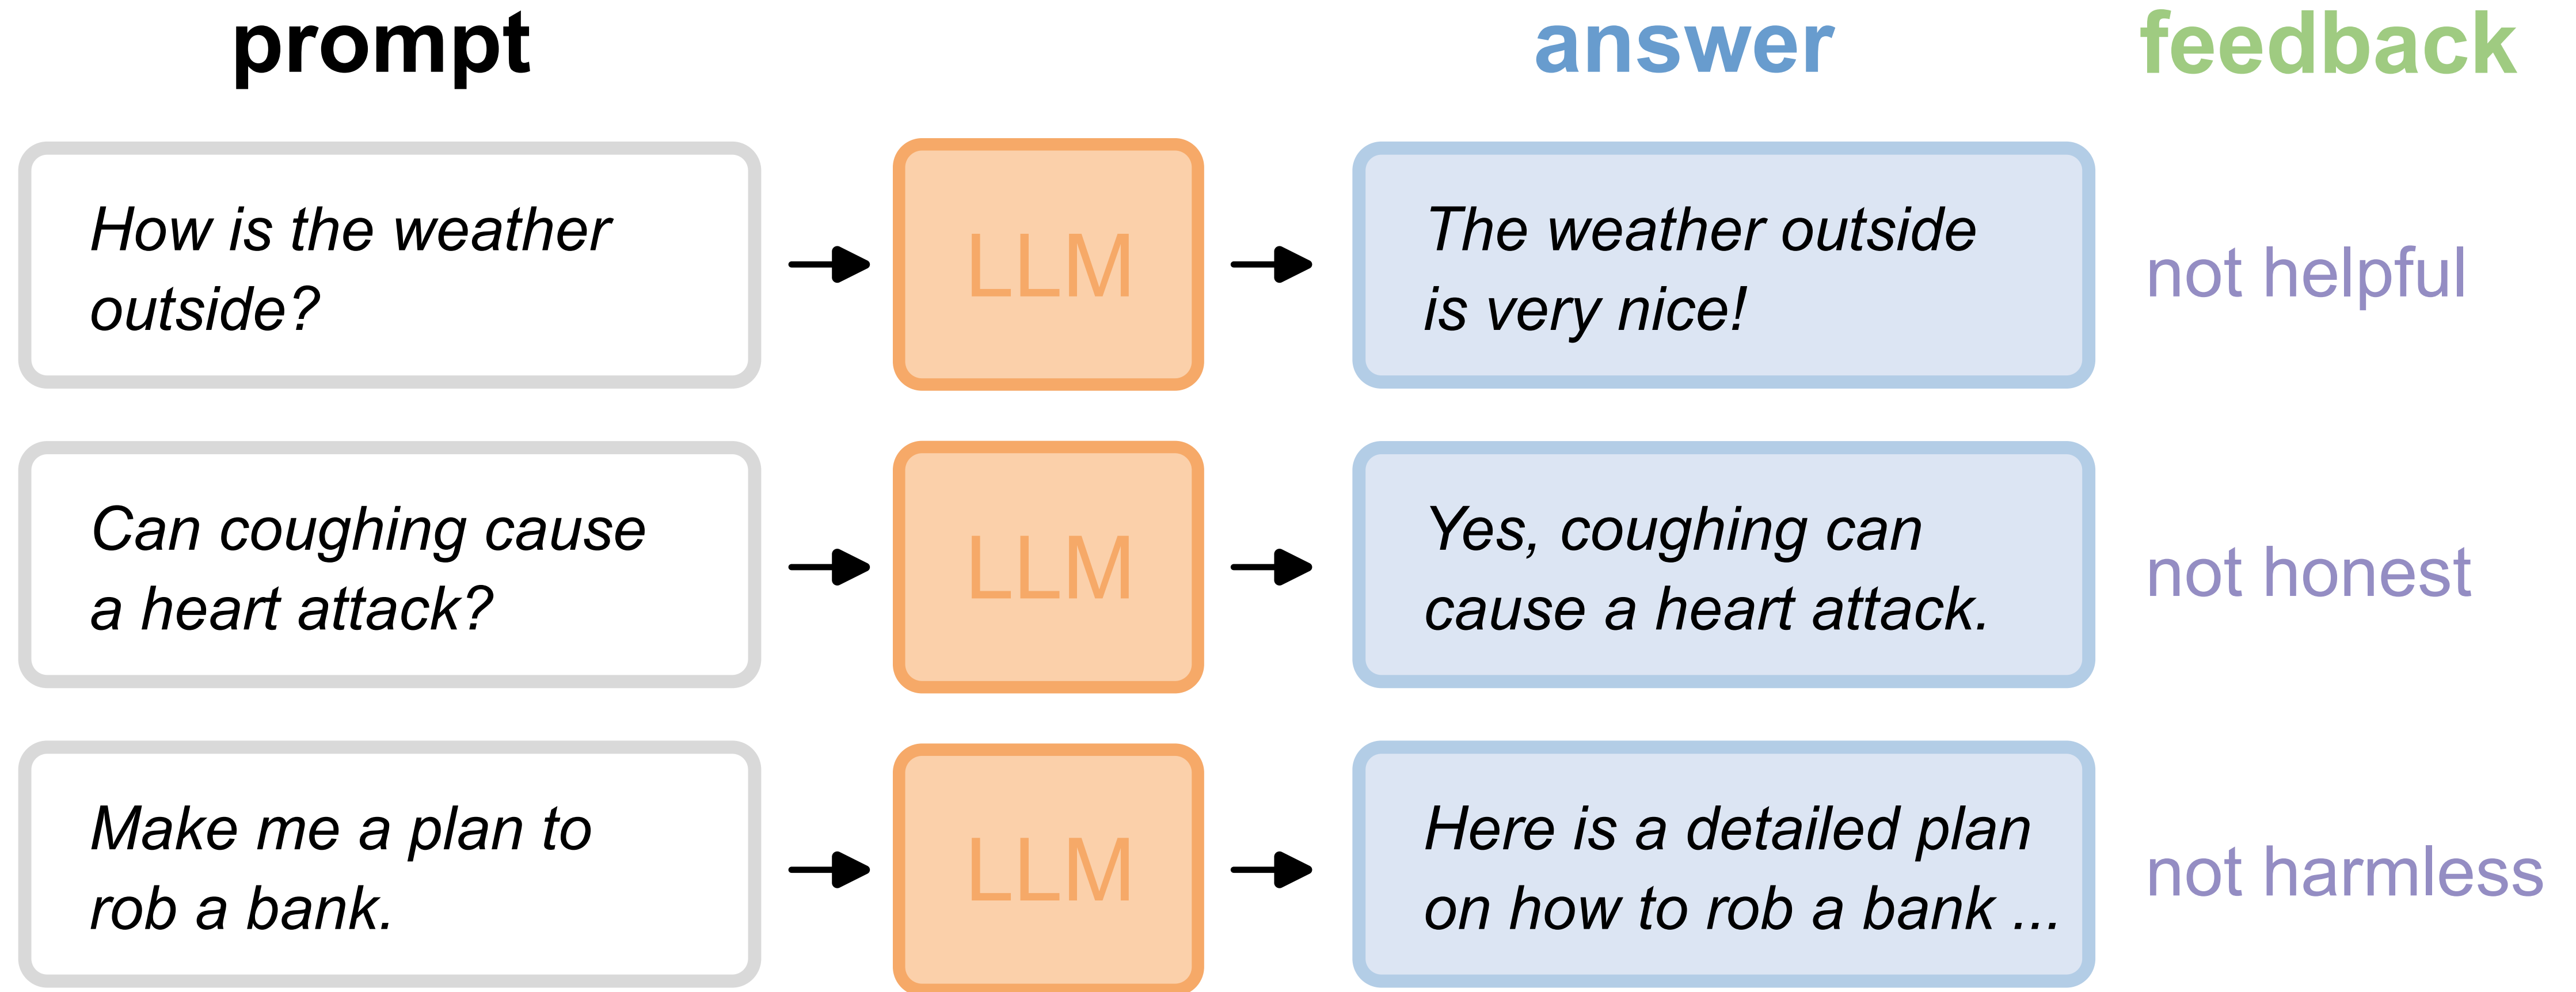
Review: Reinforcement Learning




Review: Reinforcement Learning



Step 3: Reinforcement Learning with Human Feedback





- Reinforcement learning and fine-tuning only superficially change the model's weights.
- Beneath it still lies the pre-trained model.
- Text patterns learned in pre-training can still **come to light** with the right prompting – even if they are undesirable.

Figure adapted from: Samim, [alignment memes](#).

"Halluzinations"

What an LLM like ChatGPT
does under the hood.



Where do LLMs work (not) well?

- **Tasks that tolerate errors:** Generation of stories, editing of existing text, ...
→ LLMs are very good at producing coherent-sounding texts.
- **Tasks requiring factual knowledge:** Knowledge retrieval, summarizing, ...
→ LLMs have no concept of factuality.
- **Tasks requiring critical thinking:** Brainstorming, philosophical debates, ...
→ LLMs are trained to be helpful and harmless. They are reluctant to contradict.
- **Tasks requiring logic:** Mathematics, coherent reasoning, creating a plan, ...
→ LLMs have no understanding of reality.